

Which problem do you want to address,
what is the use case you want to tackle?

We have an info@-inbox, which receives hundreds of mails per day. We want to use NLP to respond to relevant incoming emails faster without increasing the headcount.

What is the value proposition for your company?
Or: why should your company invest into this?

If we get this done, we can leverage mail automation across the company and make better use of our scarce skilled workers + keep customers happy via fast responses.

What is the value proposition for the users?
Or: why should the users care about this?

Mail senders: they get responses much faster, as less overhead is caused internally.
Departments: less overhead in emails, more time for critical work.

Which integrations / channels are part of the use case?
Or: Where is data initially coming from, and how is it processed?

We use Outlook as our mail server, so we must be able to fetch mails from it and push information back to it. Later on, we could also look into our support systems, e.g. Zendesk.

How does the data look like that is being processed in this process?
Which languages are used? (English, German, ...)

Our emails are HTML-based, i.e. can contain complex structures. We mostly communicate in English, sometimes also in German.

In this process, which automations would need to be fully automated,
where would machine and human work together?

The categorization of incoming mails (by intent, e.g. "this is about a claim") must be fully automated. Response drafting (!) would be helpful, but no fully automated responding.

How sensitive is the data processed (e.g. PII)?
Which 3rd party services are involved?

Incoming mails can potentially contain PII, so we should check for this (also as part of the NLP process). We could integrate APIs from e.g. Azure or OpenAI directly.

Where and how are we going to develop the solution?
Where and how are we going to deploy the solution?

We're going to develop the platform with Kern AI, a provider we found useful. Most likely, we're also going to deploy it with them, but maybe we'll do it on-prem.

How can we measure the validity of results?
How can we benchmark the solution?

We should have a human benchmark, i.e. check how often intents are correctly predicted. Also, we can set up some rules if the AI isn't sure enough to let a human look into the case. Maybe a 90% accuracy is what we want to aim for.

Which internal departments and teams are involved in this use case?
(Developer, process owner, domain experts, decision maker)

We need someone from the software development team to help us build a low-code solution. Also, we need the departments that are involved in the info@-inbox, and the IT.

Which external parties are involved?
(Customers, agencies)

We're going to look into this with Kern AI. They are experts in their field and have a ISO27001-certified and GDPR-conform low-code solution.

How is the cost structure of the solution we aim to build?
What are the initial setup costs and what is recurring?

We've got a budget of 20,000€ for the setup. The platform fee for our requirements is roughly 999€ monthly, but we'll start with the 149€ package in the first phase.

Value Proposition

Governance and metrics

Stakeholders and costs

Process and implementation

